

Stop calling Knative Serverless!

Doug Davis

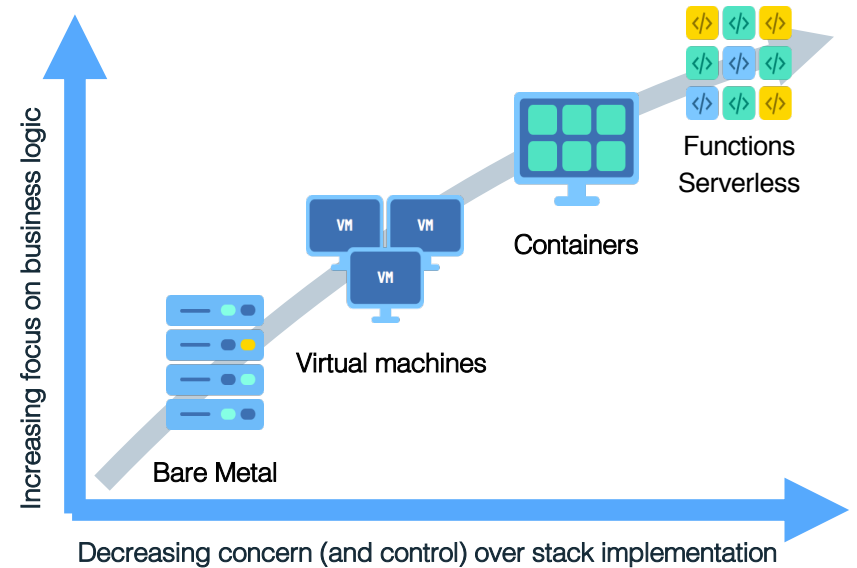
IBM - STSM - OM Knative

dug@us.ibm.com | @duginabox



Why Cloud Native?

- Break-up the monolith
 - Better resource utilization
 - Reduced costs
- Abstraction of infrastructure
 - Devs focus on code not infrastructure
 - Faster time to market = \$



Choices of infrastructure ...

- Platform as a Service
- Containers as a Service
- Functions as a Service/Serverless

e.g. CloudFoundry

e.g. Kubernetes

e.g. OpenWhisk

Which do you use?

Characteristics of Platform as a Service

	PaaS
Simplified UX	✓
Containers	✓
Micro-services / small-ish tasks	✓
Stateless	✓
Endpoint + Load Balancing	✓
Build	✓
Pay per usage (public cloud)	✓
On-demand infrastructure - auto-scaling	✓
Access to advanced features (eg. net/vol)	

Characteristics of C/F as a Service

	PaaS	CaaS	Fn/Srvr
Simplified UX	✓		✓
Containers	✓	✓	✓
Micro-services / small-ish tasks	✓	✓	✓
Stateless	✓	✓	✓
Endpoint + Load Balancing	✓	DIY	✓
Build	✓		✓
Pay per usage (public cloud)	✓	✓	✓
On-demand infrastructure - auto-scaling	✓	DIY	✓
Access to advanced features (eg. net/vol)		✓	
Event driven/tooling			✓
Scale to zero			✓
Async invocations			✓
Non-restrictive task execution times	✓	✓	
Non-restrictive resource usage (eg. mem)	✓	✓	

**Why be
forced to
choose?**



The need for something different...



Kubernetes: Most popular Container Management project today.
Deploy and manage your containerized applications.

Resources


Containers
Pods
Replica Sets
Deployments
Services
Endpoints
Secrets
Networks
Volumes/PV/PVC
Ingress/LBs

Tooling & Other "stuff"

yaml
Spec vs Status
helm
kubectl
Istio
...
Blue/green deployments

The promise of Cloud Computing:


- Higher level abstractions
- Management of infrastructure for me
 - "Decreasing concern (and control) over infrastructure implementation"



I just want to deploy my code!

Introducing Knative


- An opinionated and simplified view of application/container management
- Allowing developers to focus on coding
 - Leverage abstractions & features from PaaS, Serverless
- Built as an extension to Kubernetes
 - Still Kubernetes under the covers
 - Access to full Kubernetes if needed
 - Integrates with the non-Knative workloads
- Building blocks on which Cloud Providers can build a platform



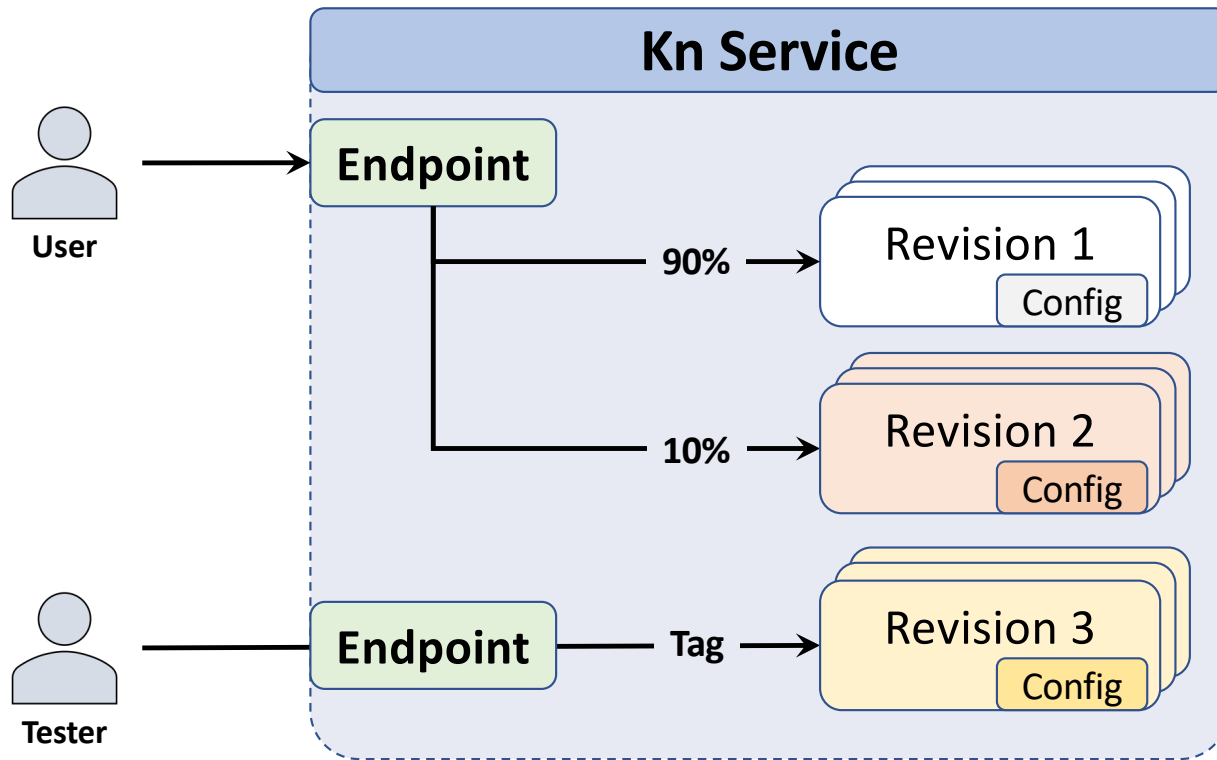
Knative

2 Main Components

- **Serving** is the runtime component
 - Host your application as K8s pods
- **Eventing** contains tools for managing events
 - Between loosely coupled services




Knative Serving - The Model



And Knative manages all of these resources for you!

- Deploy app as pod/revision
 - Revision specific config
 - E.g. image, env vars, scale
- Networking auto-setup
- Revisions are scaled up/down
 - Based on load
 - Even down to zero
- Updates create Revisions
 - Auto-migration to new
- Traffic splitting based on %
- Dedicated URLs to Revisions



Knative - Demo - The Core



Knative - Demo - The Core - Summary

- Deploy Service with just "name" and "image"
 - Access it via HTTP and HTTPS
 - Auto-scaled up and down (to zero) based on incoming load
 - Rolling upgrade to new version of Service
 - Traffic split (e.g. 50/50) between versions
 - Dedicated URL to specific version
 - Exclude specific version from traffic router
-
- What would it take to do all of this via vanilla Kubernetes?

Knative Serving - Things to know

- Container Images MUST run HTTP servers
- Multi-threaded model - but configurable
- Configuration options
 - Container Concurrency
 - Min / Max / Target Scale
 - Scaling: requests vs cpu
- Simplified resource model

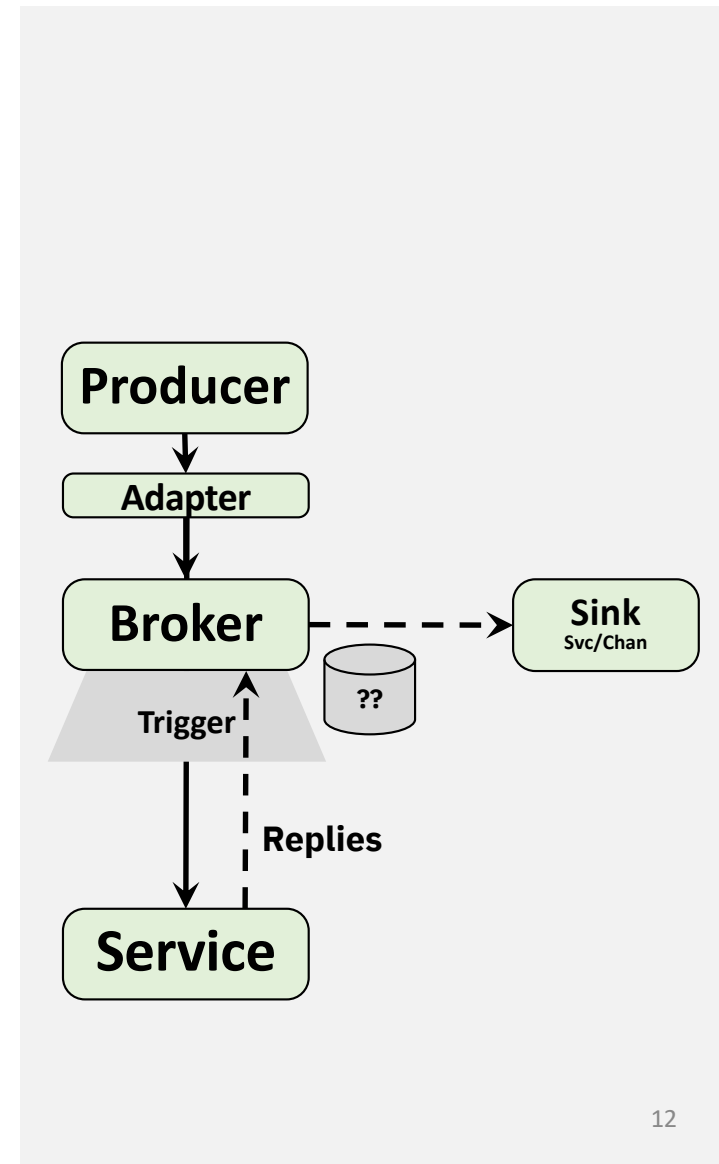
```
apiVersion: serving.knative.dev/v1beta1
kind: Service
metadata:
  name: echo
spec:
  template:
    spec:
      containers:
        - image: duglin/echo
```

Knative Eventing - Core

- **Eventing Primitives**

Manage the coordination/delivery of events


- **Event Source** - connects Event Producer to "sink"
 - Create the subscription for you
 - Often, creates an Adapter (KnService) to receive the events and convert them into CloudEvents
- **Broker** - a receiver of events
 - E.g. a queue
- **Trigger** - (subscription) ask for events from a Broker
 - Filters - to subset the stream of events
Often based on CloudEvent's metadata



Knative Eventing - Other Features

- Sequence

- Ordered set of "Sinks"
- Final "Reply"



- Choice

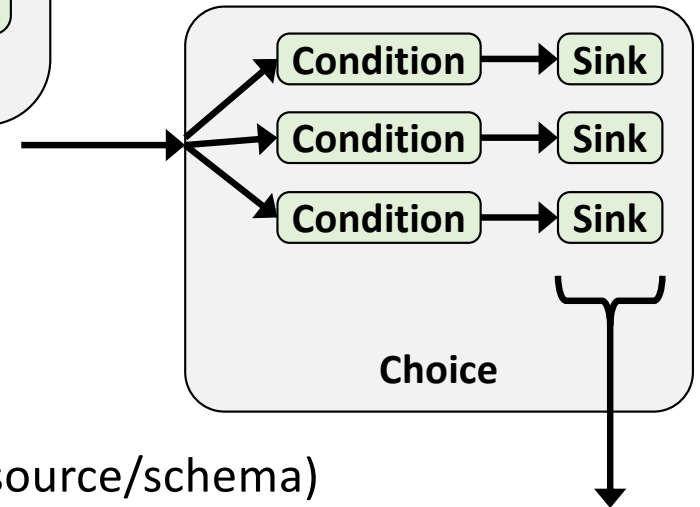
- Fan-out with filters/conditions


- Event Registry

- Collection of EventTypes (Kn Broker, CloudEvent type/source/schema)

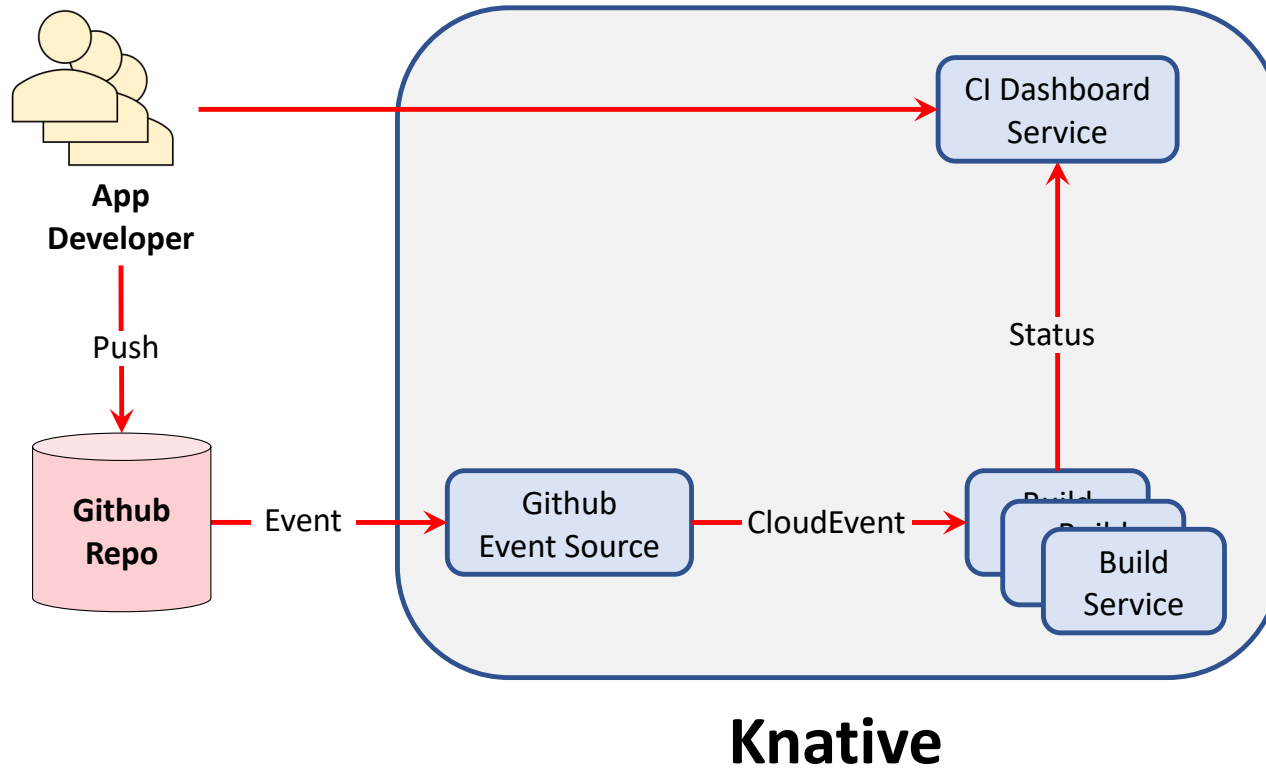
- CloudEvents - <https://cloudevents.io>

- Regardless of Producer, Transport or Format - there's consistent metadata






Knative - Demo - Build Service



- Deploy CI Dashboard KnSvc
- Deploy a "Build" KnSvc
 - Concurrency = 1
- Create a Github Event Source
 - Notified of new Commits
 - Sink = Build KnSvc
- Dev commits & checks status
- Build Service scales
- Do it !

Characteristics of * as a Service

	PaaS	CaaS	Fn/Srvr	Knative
Simplified UX	✓		✓	✓
Containers	✓	✓	✓	✓
Micro-services / small-ish tasks	✓	✓	✓	✓
Stateless	✓	✓	✓	✓
Endpoint + Load Balancing	✓	DIY	✓	✓
Build	✓		✓	WIP
Pay per usage (public cloud)	✓	✓	✓	✓
On-demand infrastructure - auto-scaling	✓	DIY	✓	✓
Access to advanced features (eg. net/vol)		✓		WIP
Event driven/tooling			✓	✓
Scale to zero			✓	✓
Async invocations			✓	WIP
Non-restrictive task execution times	✓	✓		WIP
Non-restrictive resource usage (eg. mem)	✓	✓		✓



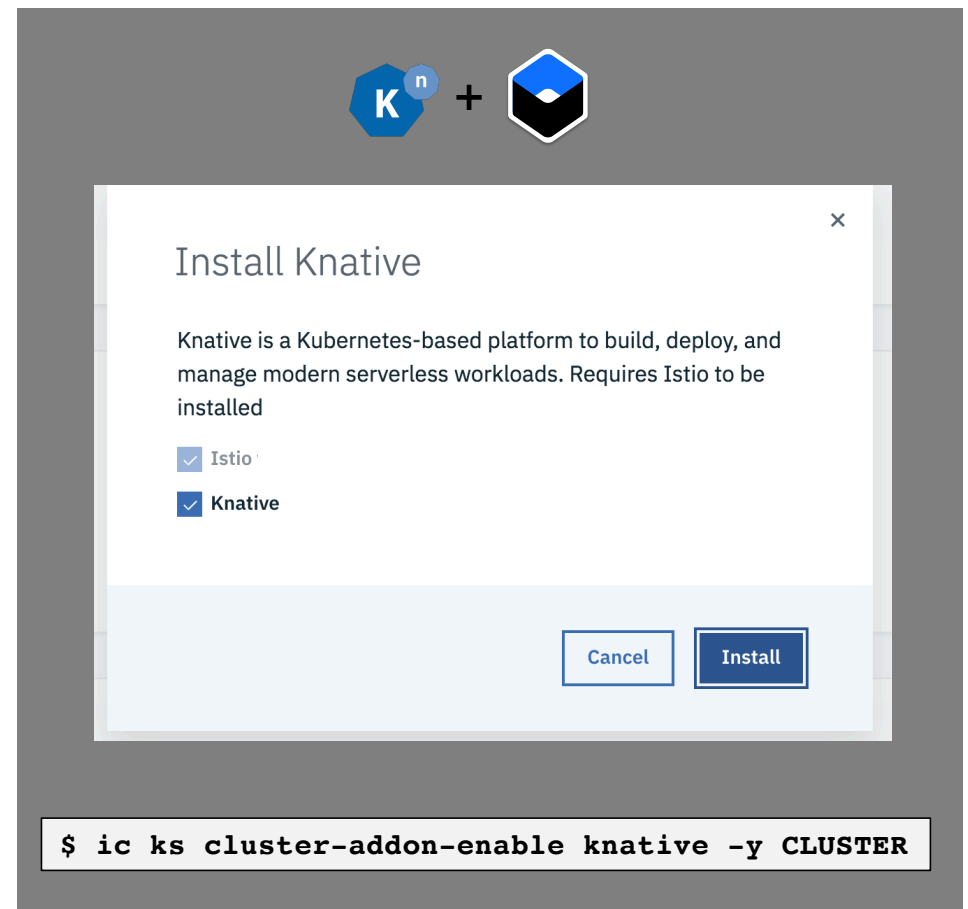
Knative - In Summary

- No longer have to choose which *aaS
- Get the best of
 - CaaS, PaaS, FaaS & Serverless
 - Simplified Kubernetes Experience
- Without losing the option of full Kubernetes when needed
- Your decisions: architecture of your application / containers
 - Size of your containers
 - Boundaries between them
 - Which bits to scale and when



IBM & Knative

- IBM involved in all parts of Knative
- Managed Knative add-on for IBM Cloud Kubernetes Service (IKS) - "experimental"
- One click install of Knative into your cluster
 - Includes Istio
- Updates to Knative are managed
- <https://ibm.com/iks>



Thank You!

Doug Davis
dug@us.ibm.com | @duginabox